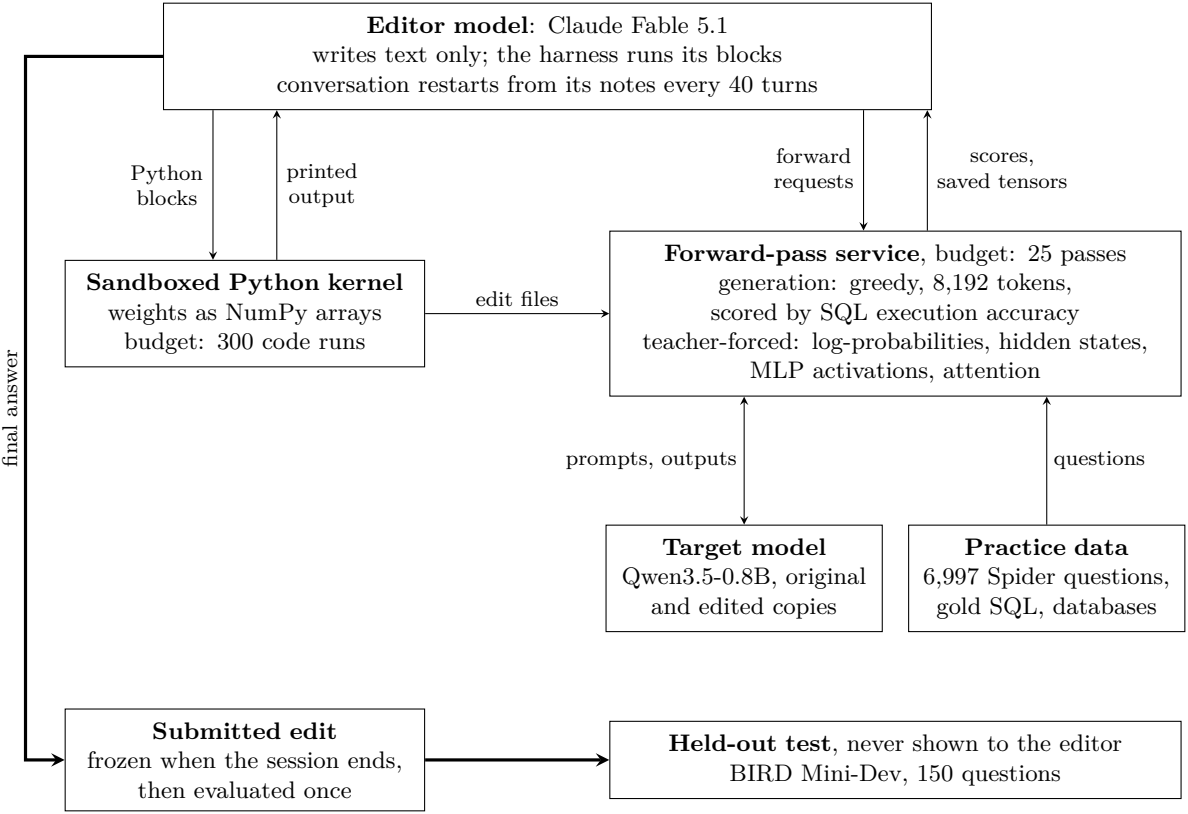
Editor model: Claude Fable 5.1
writes text only; the harness runs its blocks
conversation restarts from its notes every 40 turns
Python blocks
printed output
forward requests
scores, saved tensors
Sandboxed Python kernel
weights as NumPy arrays
budget: 300 code runs
edit files
Forward-pass service, budget: 25 passes
generation: greedy, 8,192 tokens,
scored by SQL execution accuracy
teacher-forced: log-probabilities, hidden states,
MLP activations, attention
final answer
prompts, outputs
questions
Target model
Qwen3.5-0.8B, original
and edited copies
Practice data
6,997 Spider questions,
gold SQL, databases
Submitted edit
frozen when the session ends,
then evaluated once
Held-out test, never shown to the editor
BIRD Mini-Dev, 150 questions